bulk-RNAseq with edgeR-Differentially Gene Expression (DEG) Analysis Notes

Dinçer Kılıç

2024-15-01

Following this **guide.**

1.4 Quick start

edgeR offers many variants on analyses. The glm approach is more popular than the classic approach as it offers great flexibilities. There are two testing methods under the glm framework: likelihood ratio tests and quasi-likelihood F-tests. The quasi-likelihood method is highly recommended for differential expression analyses of bulk RNA-seq data as it gives stricter error rate control by accounting for the uncertainty in dispersion estimation. The likelihood ratio test can be useful in some special cases such as single cell RNA-seq and datasets with no replicates. The details of these methods are described in Chapter 2.

A typical edgeR analysis might look like the following. Here we assume there are four RNASeq libraries in two groups, and the counts are stored in a tab-delimited text file, with gene symbols in a column called Symbol.

```
x <- read.delim("TableOfCounts.txt",row.names="Symbol")
group <- factor(c(1,1,2,2))
y <- DGEList(counts=x,group=group)
keep <- filterByExpr(y)
y <- y[keep,,keep.lib.sizes=FALSE]
y <- normLibSizes(y)
design <- model.matrix(~group)
y <- estimateDisp(y,design)
```

To perform quasi-likelihood F-tests:

```
fit <- glmQLFit(y,design)
qlf <- glmQLFTest(fit,coef=2)
topTags(qlf)
```

To perform likelihood ratio tests:

```
fit <- glmFit(y,design)
lrt <- glmLRT(fit,coef=2)
topTags(lrt)
```

2.2 Aligning reads to a genome

The first step in an RNA-seq analysis is usually to align the raw sequence reads to a reference genome, although there are many variations on this process.

We find the Subread-featureCounts pipeline [25, 26] to be very fast and effective for this purpose, but the STAR-featureCounts, STAR-htseq and Bowtie-TopHat-htseq pipelines are also popular. Subread and featureCounts are particularly convenient because it they implemented in the Bioconductor R package Rsubread [27].

2.3 Producing a table of read counts

edgeR works on a table of read counts, with rows corresponding to genes and columns to independent libraries. The counts represent the total number of reads aligning to each gene (or other genomic locus).

Such counts can be produced from aligned reads by a variety of short read software tools. We find the featureCounts function of the Rsubread package [26, 27] to be particularly effective and convenient, but other tools are available such as findOverlaps in the GenomicRanges package or the Python software htseq-counts.

Reads can be counted in a number of ways. When conducting gene-level analyses, the counts could be for reads mapping anywhere in the genomic span of the gene or the counts could be for exons only. We usually count reads that overlap any exon for the given gene, including the UTR as part of the first exon [26].

For data from pooled shRNA-seq or CRISPR-Cas9 genetic screens, the processAmplicons function [11] can be used to obtain counts directly from FASTQ files. Note that edgeR is designed to work with actual read counts. We not recommend that predicted transcript abundances are input the edgeR in place of actual counts.

2.4 Reading the counts from a file

If the table of counts has been written to a file, then the first step in any analysis will usually be to read these counts into an R session.

If the count data is contained in a single tab-delimited or comma-separated text file with multiple columns, one for each sample, then the simplest method is usually to read the file into R using one of the standard R read functions such as read.delim.

If the counts for different samples are stored in separate files, then the files have to be read separately and collated together. The edgeR function readDGE is provided to do this. Files need to contain two columns, one for the counts and one for a gene identifier.

2.5 Pseudoalignment and quasi-mapping

The kallisto and Salmon software tools align sequence reads to the transcriptome instead of the reference genome and produce estimated counts per transcript. Output from either tool can be input to edgeR via the tximport package, which produces gene-level estimated counts and an associated edgeR offset matrix. Alternatively, kallisto or Salmon output can be read directly into edgeR using the catchSalmon and catchKallisto functions if the intention is to conduct a transcript-level analysis.

2.6 The DGEList data class

edgeR stores data in a simple list-based data object called a DGEList. This type of object is easy to use because it can be manipulated like any list in R. The function readDGE makes a DGEList object directly. If the table of counts is already available as a matrix or a data.frame, x say, then a DGEList object can be made by

```
y <- DGEList(counts=x)
```

A grouping factor can be added at the same time:

```
group <- c(1,1,2,2)
y <- DGEList(counts=x, group=group)
```

The main components of an DGEList object are a matrix counts containing the integer counts, a data.frame samples containing information about the samples or libraries, and a optional data.frame genes containing annotation for the genes or genomic features. The data.frame samples contains a column lib.size for the library size or sequencing depth for each sample. If not specified by the user, the library sizes will be computed from the column sums of the counts. For classic edgeR the data.frame samples must also contain a column group, identifying the group membership of each sample.

2.7 Filtering

Genes with very low counts across all libraries provide little evidence for differential expression. In the biological point of view, a gene must be expressed at some minimal level before it is likely to be translated into a protein or to be biologically important. In addition, the pronounced discreteness of these counts interferes with some of the statistical approximations that are used later in the pipeline. These genes should be filtered out prior to further analysis.

As a rule of thumb, genes are dropped if they can't possibly be expressed in all the samples for any of the conditions. Users can set their own definition of genes being expressed. Usually a gene is required to have a count of 5-10 in a library to be considered expressed in that library. Users should also filter with count-per-million (CPM) rather than filtering on the counts directly, as the latter does not account for differences in library sizes between samples.

Here is a simple example. Suppose the sample information of a DGEList object y is shown as follows:

```
 y$samples
        group lib.size norm.factors

Sample1 1     10880519 1
Sample2 1     9314747  1
Sample3 1     11959792 1
Sample4 2     7460595  1
Sample5 2     6714958  1
```

We filter out lowly expressed genes using the following commands:

```
keep <- filterByExpr(y)
y <- y[keep, , keep.lib.sizes=FALSE]
```

The filterByExpr function keeps rows that have worthwhile counts in a minumum number of samples (two samples in this case because the smallest group size is two). The function accesses the group factor contained in y in order to compute the minimum group size, but the filtering is performed independently of which sample belongs to which group so that no bias is introduced. It is recommended to recalculate the library sizes of the DGEList object after the filtering, although the downstream analysis is robust to whether this is done or not.

The group factor or the experimental design matrix can also be given directly to the filterBy Expr function by

```
keep <- filterByExpr(y, group=group)
```

if not already set in the DGEList object. More generally, filterByExpr can be used with any design matrix:

```
keep <- filterByExpr(y, design)
```

The filtering should be based on the grouping factors or treatment factors that will be involved in the differential expression teststested for, rather than on blocking variables that are not of scientific interest in

themselves. For example, consider a paired comparison experiment in which the same treatment regimes applied to each of a number of subjects or patients:

```
design <- model.matrix(~ Patient + Treatment)
```

In this design, Patient is included in the design matrix to correct for baseline differences between the Patients, but we will not be testing for differential expression between the Patients. The filtering should therefore be based soley Treatment rather than on Patient, i.e

```
keep <- filterByExpr(y, group=Treatment)
```

rather than

```
keep <- filterByExpr(y, design)
```

2.8 Normalization

```
y <- normLibSizes(y)
y$samples

        group lib.size norm.factors

Sample1 1     10880519 1.17
Sample2 1      9314747 0.86
Sample3 1     11959792 1.32
Sample4 2      7460595 0.91
Sample5 2      6714958 0.83
```

2.9 Negative binomial models

2.10 The classic edgeR pipeline: pairwise comparisons between two or more groups

2.10.1 Estimating dispersions

To estimate common dispersion and tagwise dispersions in one run (recommended):

```
y <- estimateDisp(y)
```

Alternatively, to estimate common dispersion:

```
y <- estimateCommonDisp(y)
```

Then to estimate tagwise dispersions:

```
y <- estimateTagwiseDisp(y)
```

```
et <- exactTest(y)
topTags(et)
```

2.11 More complex experiments (glm functionality)

2.11.2 Estimating dispersions

```
y <- estimateDisp(y, design)
```

Alternatively, one can use the following calling sequence to estimate them one by one. To estimate common dispersion:

```
y <- estimateGLMCommonDisp(y, design)
```

To estimate trended dispersions:

```
y <- estimateGLMTrendedDisp(y, design)
```

To estimate tagwise dispersions:

```
y <- estimateGLMTagwiseDisp(y, design)
```

2.11.3 Testing for DE genes

As a brief example, consider a situation in which are three treatment groups, each with two replicates, and the researcher wants to make pairwise comparisons between them. A QL model representing the study design can be fitted to the data with commands such as:

```
group <- factor(c(1,1,2,2,3,3))
design <- model.matrix(~group)
fit <- glmQLFit(y, design)
```

The fit has three parameters. The first is the baseline level of group 1. The second and third are the 2 vs 1 and 3 vs 1 differences.

To compare 2 vs 1:

```
 qlf.2vs1 <- glmQLFTest(fit, coef=2)
topTags(qlf.2vs1)
```

To compare 3 vs 1:

```
qlf.3vs1 <- glmQLFTest(fit, coef=3)
```

To compare 3 vs 2:

```
qlf.3vs2 <- glmQLFTest(fit, contrast=c(0,-1,1))
```

To find genes different between any of the three groups:

```
qlf <- glmQLFTest(fit, coef=2:3)
topTags(qlf)
```

Alternatively, one can perform likelihood ratio test to test for differential expression. The testing can be done by using the functions glmFit() and glmLRT(). To apply the likelihood ratio test to the above example and compare 2 vs 1:

```
fit <- glmFit(y, design)
lrt.2vs1 <- glmLRT(fit, coef=2)
topTags(lrt.2vs1)
```